

DentaCoPilot: An LLM-Augmented Next-Procedure Recommender for General Dentistry, Designed for Dentist Augmentation

Carson Conception Rodrigues^{a,*}, Steffie Dione Rebello^b

^a*CELABE, Research Division*

^b*KLE Vishwanath Katti Institute of Dental Sciences, Belgaum, India*

Abstract

Background. Commercial dental artificial intelligence in 2026 is overwhelmingly diagnostic: caries, calculus, periapical, and bone-level detection on radiographs. The clinically harder question that follows every diagnosis — given a patient’s chart and most recent procedure, what should the dentist do next — remains unsolved at general-dentistry scale. The closest published system, MultiTP (Chen et al., 2024), is a CNN-RNN restricted to partial-edentulism cases and provides neither calibrated uncertainty, structured rationale, nor an evaluation that treats the model as decision support rather than as an autonomous classifier. **Methods.** We introduce *DentaCoPilot*, a recommender that, given a structured chart, returns (i) a calibrated top- K probability distribution over Current Dental Terminology (CDT) codes for the next procedure, (ii) a verbalised confidence label, (iii) an explicit *abstain* flag when context is insufficient, and (iv) a chart-grounded rationale. We compare four classical baselines (frequency bigram, TF-IDF + logistic regression, XGBoost, MultiTP-style CNN-RNN) and six large-language-model (LLM) variants (Claude Haiku, Sonnet + chain-of-thought, Sonnet + retrieval, Opus + chain-of-thought, Sonnet + classical prior, Opus + classical prior) on a synthetic chart corpus of 500 patients (1,284 test examples). All LLM inference is routed through the local Anthropic Claude Code CLI; every call is logged for full audit. **Results.** On apples-to-apples evaluation, classical baselines reach 0.567 top-1 / 0.967 top-5; pure LLM variants trail at 0.267–0.467 top-1. Prompt-conditioning a Sonnet LLM on the classical baseline’s top-10 candidates (M5) closes the gap: top-5 rises from 0.733 (pure Sonnet + chain-of-thought) to 0.933, matching classical baselines, while preserving rationale and abstention. Increasing

the LLM backbone from Sonnet to Opus does not improve accuracy with or without priming. Calibration via temperature scaling and coverage-risk analysis is reported for the baselines. **Conclusion.** Prompt-conditioning a small LLM on a classical baseline’s top-K is the most cost-effective LLM design we tested for next-procedure recommendation, and the design preserves the augmentation features that distinguish the system from an autonomous classifier. A pre-registered clinician-in-the-loop evaluation at the KLE Vishwanath Katti Institute of Dental Sciences (Belgaum, India) and a real-data evaluation on the multi-institutional BigMouth dental data repository are the next stage of work.

Keywords: dental artificial intelligence, large language models, clinical decision support, procedure prediction, human-AI collaboration, calibrated abstention

1. Introduction

Commercial dental AI in 2026 is overwhelmingly diagnostic. The most recent FDA-clearance survey (Shujaat, 2026) catalogues 29 cleared products from 13 companies (Overjet, Pearl, VideaHealth, Dexis, Perimetrics, and others) targeting single-image classification problems — caries detection, calculus identification, periapical radiolucency, and periodontal bone-level measurement.¹ Peer-reviewed evidence for clinical benefit exists for the caries-detection variant in particular: in a cluster-randomised cross-over trial of 22 dentists assessing bitewings, Mertens et al. (2021) reported that AI-assisted review raised the mean diagnostic-accuracy area under the ROC from 0.85 to 0.89 ($\Delta\text{AUC} = +0.04$, $p < 0.05$); a follow-up trial confirmed the gain was cost-effective at practice scale (Schwendicke et al., 2022). These are valuable deployments. They are also, individually, single-decision tools. They do not answer the clinically harder question that follows every diagnosis: *what should the dentist do next?*

Sequencing a treatment plan from a chart is a different problem. A two-surface posterior composite on tooth 19 leads to one set of likely continuations; the same composite on a tooth with a prior endodontic history

*Corresponding author.

Email addresses: carson@celabe.com (Carson Conception Rodrigues), steffie1619@gmail.com (Steffie Dione Rebello)

¹See American Dental Association (2024) for the procedure code set used throughout this paper.

Preprint submitted to Journal of Dental Research

May 8, 2026

19 leads to another; in a periodontitis patient with diabetes, neither sequence is
 20 appropriate. The closest published system that operates at this layer, Mul-
 21 tiTP (Chen et al., 2024), is a CNN-RNN restricted to partial-edentulism
 22 cases with a fixed five-treatment output sequence. It does not surface cali-
 23 brated uncertainty, does not abstain when the chart context is insufficient,
 24 and was not evaluated as a clinician-augmenting tool. The largest related
 25 effort in periodontitis-risk prediction (Patel et al., 2026) is a risk model, not
 26 a procedure recommender.

27 We propose *DentaCoPilot*: an LLM-augmented next-procedure recom-
 28 mender designed from the start as decision support for dentists. The system
 29 takes a structured chart and the most recently completed procedure, and
 30 emits four signals together: (i) a top- K probability distribution over Current
 31 Dental Terminology codes for the next procedure; (ii) verbalised confidence
 32 in {low, medium, high}; (iii) an explicit **abstain** flag when the chart con-
 33 text is genuinely insufficient; (iv) a chart-grounded rationale citing specific
 34 prior procedures and findings. The model is constrained-decoded to the
 35 CDT 2024 vocabulary and calibrated via temperature scaling. Critically,
 36 the contribution lies less in raw accuracy than in the framing: we evaluate
 37 the system as a consult that augments the dentist, not as a classifier that
 38 replaces the dentist’s judgement.

39 We test two hypotheses on a synthetic chart corpus of 500 patients (§3.2):
 40 **(H1)** a calibrated LLM-based recommender can match classical-baseline ac-
 41 curacy on next-procedure prediction while preserving rationale and abstention;
 42 **(H2)** injecting a classical baseline’s top- K candidates into the LLM
 43 prompt is the most cost-effective LLM design for achieving (H1). We evalu-
 44 ate four classical baselines, six LLM variants spanning Claude Haiku, Son-
 45 net, and Opus with chain-of-thought, retrieval, and prior-conditioning, and
 46 report calibration and coverage–risk analyses. A pre-registered clinician-in-
 47 the-loop evaluation at the KLE Vishwanath Katti Institute of Dental Sci-
 48 ences (Belgaum, India) and a real-data evaluation on the multi-institutional
 49 BigMouth dental data repository (Walji et al., 2014, 2022) are the planned
 50 next stage of work.

51 All LLM inference in this work is routed through the local Anthropic
 52 Claude Code CLI (Anthropic, 2026), with full per-call audit logs released
 53 alongside the code (see §3.6).

54 2. Related Work

55 *Diagnostic dental AI in commercial deployment..* Shujaat (2026) catalogues

56 the current FDA-cleared landscape: caries / calculus detection (Overjet,
57 Pearl, VideaHealth), periapical radiolucency (Overjet, Pearl), automated
58 dental charting (Overjet), millimetre-precision bone-level quantification (Over-
59 jet, Pearl), 3D cone-beam CT segmentation (Pearl), and acoustic fracture
60 detection (Perimetrics). These systems share a common shape: an image
61 goes in, a finding comes out. None plans treatment.

62 *Procedure sequencing..* MultiTP (Chen et al., 2024) is the closest published
63 system to what we propose: a CNN-RNN with attention that predicts up to
64 five sequential treatments for partial-edentulism cases, achieving an AUC
65 of 0.91 over five plans. It is an EDR-only model with no calibration, no
66 abstention, and no clinician-augmentation evaluation. We reimplement a
67 minimal MultiTP-style architecture as our B3 comparator. In a related
68 specialty domain, Li and Wang (2025) use multi-task reinforcement learn-
69 ing with explainable AI to plan orthodontic-orthognathic treatment, but
70 the task is sufficiently different (continuous decision-making over a longer
71 planning horizon, narrow specialty) that the contributions do not transfer
72 directly to general dentistry.

73 *Risk prediction in dentistry..* Patel et al. (2026) model periodontitis risk
74 on 20,946 adult patients with linked EHR-EDR data, demonstrating that
75 procedure-level structured data can support clinical prediction; the same
76 group has also shown that adding social determinants of health to such
77 models meaningfully changes risk calibration (Patel et al., 2025). Varghese
78 et al. (2025) pair Cariogram-based caries-risk profiling with parent-oriented
79 educational mobile messages in a randomised trial, an example of coupling
80 structured risk modelling with a patient-facing digital intervention. All
81 three are risk models, not procedure recommenders; we position our work
82 as the procedure-recommendation analogue.

83 *LLMs in dentistry..* Farhadi Nia et al. (2024) surveys early LLM applica-
84 tions in dental practice, focusing on documentation, summarisation, and
85 patient-facing chat. The review explicitly flags that LLMs “present clean,
86 fluent statements without communicating uncertainty,” which is the fail-
87 ure mode this paper directly addresses through calibration and abstention.
88 Broader medicine-LLM work (Singhal et al., 2023; Nori et al., 2023) shows
89 that capable models reach physician-level accuracy on knowledge bench-
90 marks, but Omiye et al. (2024) demonstrates that the same models can
91 propagate dangerous race-based clinical claims — which motivates our hard

requirement for chart-grounded rationale and explicit abstention. Schwendicke et al. (2020) surveys the chances and challenges of dental AI broadly. To our knowledge, no prior work has built and evaluated an LLM as a treatment-sequence recommender, and no prior work has done so under a clinician-augmentation rather than autonomy framing.

Calibration and abstention in clinical AI. Our calibration approach follows Guo et al. (2017) (temperature scaling) with Platt (2000) as the precursor. Coverage– risk and selective prediction are standard tools in safety-critical ML; we apply them to the procedure-recommendation setting and pre-register the threshold-selection rule before any test-split unblinding.

Datasets. BigMouth (Walji et al., 2014, 2022) is a multi-institutional dental EHR repository with 6M+ patients across multiple dental schools, accessed under a project- review-committee proposal. Open imaging datasets relevant for chart- context grounding include DENTEX (Hamamci et al., 2023), MMDental (Wang et al., 2025), and DenPAR (Rasnayaka et al., 2025).

3. Methods

3.1. Task formulation

We formulate next-procedure prediction as a single-step classification over the Current Dental Terminology code set (American Dental Association, 2024; Kalenderian et al., 2018). A patient chart is a tuple $C = (D, M, H)$ where D holds demographics (age band, sex), M is a list of medical-history flags, and $H = (e_1, e_2, \dots, e_T)$ is a time-ordered sequence of procedure events. Each event $e_t = (\tau_t, c_t, k_t)$ carries a date τ_t , a CDT code $c_t \in \mathcal{V}$, and an optional Universal-numbered tooth identifier $k_t \in \{1, \dots, 32\} \cup \{\perp\}$ following standard ADA conventions (Wadia, 2017).

Given a chart truncated at index t (i.e. the dentist has just performed e_t), the model outputs:

1. A top- K probability distribution $\hat{p}(c_{t+1} \mid C_{1:t}) \in \Delta^{|\mathcal{V}|}$ over the next procedure code, returned as the top- K codes ranked by probability.
2. A confidence label $\hat{q} \in \{\text{low}, \text{medium}, \text{high}\}$, a verbalised self-rating (Lin et al., 2022; Kadavath et al., 2022; Xiong et al., 2024).
3. An abstain flag $a \in \{0, 1\}$ which the system sets to 1 when the chart context is judged insufficient.
4. A short rationale r grounded in cited fields of the chart $C_{1:t}$ (used for evaluation, not for the prediction itself).

127 We expand each chart of length L into $L - 1$ prediction examples by
 128 truncating at every prefix index $t \in \{1, \dots, L - 1\}$. Train, dev, and test
 129 splits are constructed at the *patient* level (no chart appears in two splits),
 130 following standard EHR-modeling practice (Rajkomar et al., 2018; Park
 131 et al., 2024).

132 We report **top- K accuracy**: the fraction of test examples for which the
 133 gold next-procedure code appears among the model’s K highest-probability
 134 predictions. Top-1 corresponds to exact CDT match; top-3 and top-5 cap-
 135 ture clinically reasonable shortlists for a recommender (a dentist seeing a
 136 small candidate list with calibrated probabilities is the deployment mode we
 137 have in mind). We also report top-1 by CDT category and, where calibrated
 138 probabilities are available, Expected Calibration Error and coverage–risk
 139 under abstention (§3.4).

140 3.2. Datasets

141 *Synthetic chart corpus.* Because BigMouth approval and KLE IRB are
 142 pending, this paper’s empirical experiments use a synthetic corpus gener-
 143 ated by a hierarchical Markov-style procedure-history generator. Patients
 144 are assigned to one of five trajectory archetypes (*healthy*, *restorative-prone*,
 145 *periodontitis-prone*, *edentulous-progressing*, *ortho-adolescent*) with plausibility-
 146 based priors. Within each archetype, per-archetype transition tables encode
 147 bigram probabilities $\Pr(c_{t+1} \mid c_t, \text{archetype})$ drawn from clinical workflow
 148 knowledge and the published procedure-frequency analyses in the BigMouth
 149 corpus (Walji et al., 2014, 2022). Tooth-level state is tracked: extracted
 150 teeth cannot receive subsequent restorations; root-canal-treated teeth tend
 151 to receive crowns next; SDF-treated sites revert to monitoring (Slayton
 152 et al., 2018; Seifo et al., 2020). We generate 500 patients over a five-year
 153 timeline, expand to 7,162 prediction examples, and split 70/15/15 at the
 154 patient level (350 / 75 / 75 charts).

155 *Public imaging datasets (planned, for chart-context grounding).* Loaders
 156 are implemented and verified for DENTEX (MICCAI 2023 grand-challenge
 157 (Hamamci et al., 2023); 2,332 panoramic X-rays with caries, deep caries,
 158 periapical, and impacted-tooth labels); MMDental (Wang et al., 2025),
 159 with 660 paired CBCT volumes and structured records; and DenPAR (Ras-
 160 nayaka et al., 2025), with annotated periapical radiographs. We do not
 161 auto-download these datasets; the loaders document the manual fetch step
 162 and verify SHA-256.

163 *Real-data evaluation (forthcoming)*.. We have applied to the BigMouth Den-
164 tal Data Repository (Walji et al., 2014, 2022) for controlled access and are
165 preparing the IRB protocol for a retrospective chart pull at the KLE Vish-
166 wanath Katti Institute of Dental Sciences, Belgaum. Both are required for
167 the clinician-in-the-loop evaluation pre-registered on OSF (Munafò et al.,
168 2017; Nosek et al., 2018).

169 3.3. Models

170 We evaluate four classical baselines and six LLM variants. All operate
171 on the same patient-level splits and the same prediction examples. CDT
172 codes outside the training-set vocabulary are masked out at inference; LLM
173 responses are post-hoc constrained to the training vocabulary.

174 *Classical baselines.*

175 **B0 frequency bigram** $\hat{p}(c_{t+1} | c_t)$ computed from training data with add-
176 one smoothing. The trivial sanity floor.

177 **B1 TF-IDF + multinomial logistic regression** (Salton and Buckley, 1988;
178 Pedregosa et al., 2011). Tokens include each historical CDT code,
179 positional-from-end markers for the last three procedures, and cate-
180 gorical demographics.

181 **B2 XGBoost** (Chen and Guestrin, 2016) on engineered chart-history fea-
182 tures: counts per CDT category, has-extraction / has-RCT / has-
183 implant / has-crown booleans, time since last visit, age band one-hot,
184 sex, and a small medical-history flag set.

185 **B3 MultiTP-style CNN–RNN** (Chen et al., 2024). Minimal reimple-
186 mentation: 32-dim code embedding, 1-D CNN over the most recent
187 procedures (analogous to MultiTP’s tooth + neighbour tensor), GRU
188 over the longer sequence, attention pooling, softmax head over the
189 CDT vocabulary. $\sim 30k$ parameters, 12 epochs.

190 *LLM variants.* All LLM inference is routed through the local Anthropic
191 Claude Code CLI (§3.6), not the metered Anthropic API.

192 **M1 Claude Haiku** (claude-haiku-4-5-20251001). Structured prompt:
193 chart pretty-printed, top- K instruction, JSON-only output schema,
194 top-40 most-frequent training codes shown as a vocabulary palette.

- 195 **M2 Claude Sonnet + chain-of-thought** (`claude-sonnet-4-6`). Same
 196 prompt as M1 but with an explicit `<thinking>...</thinking>` rea-
 197 soning step before the JSON, in line with prior work showing CoT
 198 prompting improves reasoning-heavy clinical tasks (Wei et al., 2022;
 199 Nori et al., 2023).
- 200 **M3 Claude Sonnet + retrieval** (*M2* + BM25 (Robertson et al., 1995)).
 201 The top-3 BM25-retrieved cards from a twenty-card open-source clin-
 202 ical evidence corpus (paraphrased from AAE / AAP / AAOMS stan-
 203 dard practice; original ADA-copyrighted material is not reproduced)
 204 are injected into the prompt; the model is instructed to cite specific
 205 cards (`[E###]`) in its rationale.
- 206 **M4 Claude Opus + chain-of-thought** (`claude-opus-4-7`). Identical prompt
 207 and schema to M2; only the model identifier changes. Tests whether
 208 more capable backbone helps.
- 209 **M5 Sonnet + classical prior** — *the proposed forward design*. The B1
 210 (TF-IDF + LR) baseline is fit on the same training set; its top-10
 211 (code, probability) candidates are injected as a prior into the M2
 212 prompt. The LLM is instructed to re-rank, redistribute, or override.
- 213 **M6 Opus + classical prior** — M5 with Opus backbone. Tests whether
 214 priming compounds with capability.

215 3.4. Calibration and abstention

216 We apply temperature scaling (Guo et al., 2017; Platt, 2000) to the top-
 217 1 probability of each model. The temperature T is fit on a random half
 218 of the test split by minimising negative log-likelihood (Naeini et al., 2015);
 219 the rescaled top-1 probability is reported on the held-out half. We compute
 220 the Expected Calibration Error (Naeini et al., 2015; Nixon et al., 2019) on
 221 equal-width bins.

222 For abstention, we sweep a threshold τ over the calibrated top-1 proba-
 223 bility and report coverage–risk curves (El-Yaniv and Wiener, 2010; Geifman
 224 and El-Yaniv, 2017); the pre-registered abstention rule (Nosek et al., 2018)
 225 is the smallest τ achieving $\geq 95\%$ precision on covered cases at $\geq 50\%$
 226 coverage on the dev split, locked before any test-split unblinding.

3.5. Hybrid integration

We test two ways to combine classical and LLM predictions:

1. **Confidence-routed gating.** Pick the baseline’s output when its top-1 probability $\geq \tau$; otherwise the LLM’s. We sweep $\tau \in \{0, 0.2, \dots, 1.0\}$.
2. **Prompt-conditioning (M5/M6).** The classical model’s top-10 (code, probability) candidates are part of the LLM prompt; the LLM produces a single output that uses, refines, or overrides the prior.

(2) is implemented end-to-end as M5 and M6; (1) is computed post-hoc from the predictions logged by each model.

3.6. Compute infrastructure

Every LLM inference call in this paper runs through a single subprocess wrapper at `data/cc_compute.py`, which invokes the local Anthropic Claude Code CLI (Anthropic, 2026) with `--output-format json` and parses the response. The wrapper logs every call (timestamp, run id, model, prompt, response, latency, input and output token counts) to a per-run `llm_calls.jsonl`, which the repository commits as a permanent record. We deliberately do not import the Anthropic Python SDK; a continuous-integration grep guard (`check-no-anthropic` in the Makefile) prevents accidental introduction. This design (a) routes compute to the existing Claude Code subscription, (b) provides a complete audit trail every reviewer can re-read without spending compute, and (c) makes the experiment reproducible by simple replay instead of re-execution.

3.7. Reproducibility protocol

All seeds are fixed (`random.seed(42)`, `numpy.random.seed(42)`, `torch.manual_seed(42)`). Each experimental run writes `config.json` (git SHA, seed, CLI args, Python and platform version), `results.json` (per-trial predictions and gold labels), and `summary.json` (aggregated metrics). The synthetic generator, all baseline and LLM implementations, the evaluation harness, and the analysis scripts are released under the MIT licence. The pre-registration (Nosek et al., 2018) is deposited on OSF.

4. Experiments

4.1. Classical baseline accuracy at scale

We first evaluate the four classical baselines (B0–B3, defined in §3.3) on the full test split (1,284 examples). This is the larger- n regime; LLM evaluation in §4.4 is necessarily smaller.

Table 1: Synthetic-corpus baseline results ($N=500$ charts; 1,284 test examples).

Model	Top-1	Top-3	Top-5
B0 frequency bigram	0.556	0.778	0.882
B1 TF-IDF + LR	0.558	0.832	0.917
B2 XGBoost	0.507	0.815	0.918
B3 MultiTP-style CNN-RNN	0.556	0.847	0.935

The four baselines cluster tightly on top-1 (50.7–55.8%) but spread on top- $K>1$. The MultiTP-style sequence model (B3) wins both top-3 (0.847) and top-5 (0.935), suggesting that CNN+RNN attention captures useful ordering of treatment continuations beyond what a flat bigram or feature classifier can. Top-1 is harder: the bigram baseline’s strong local Markov signal nearly ties B1 (0.556 vs 0.558), and XGBoost lags both, indicating the engineered feature set under-uses recent-procedure information.

The per-category breakdown (Figure 2) reveals the opportunity for chart-aware reasoning: classical baselines achieve 75–80% top-1 on bundled / sequential procedures (diagnostic, adjunctive, orthodontic, preventive) but *collapse* on the multi-class restorative choice (1% top-1 across all three baselines), endodontic sequencing (0%, scarce in synthetic data), and prosthodontic decisions (0%). This is exactly the regime where a chart-aware LLM with calibrated abstention should help. The next subsection (§4.4) reports the LLM evaluation against these baselines.

4.2. Reproducibility (run records)

Every experimental run is recorded as described in §3.7. Per-run `config.json`, `results.json`, `summary.json`, and (for LLM-backed runs) `llm_calls.jsonl` are committed to the repository. All large language model inference in this paper runs through the local Anthropic Claude Code CLI instead of the metered Anthropic API; the wrapper at `data/cc_compute.py` is the single chokepoint.

4.3. Calibration and abstention

Following the protocol in §3.4, we post-hoc calibrate each baseline and report ECE pre- and post-temperature-scaling.

XGBoost is materially overconfident ($T=1.75$ needed; calibration does not converge in our held-out split). The bigram and TF-IDF + LR models




Figure 1: Top-K next-procedure accuracy on the apples-to-apples $n=30$ shared subsample. The four classical baselines (B0–B3) outperform the three LLM variants (M1–M3) on top-1 exact-CDT-match; the gap narrows at top-K=3 and 5. M2 (Sonnet + chain-of-thought) is the strongest LLM variant. M3 (Sonnet + retrieval) does not improve over M2 on this small synthetic test, suggesting that retrieval needs richer chart contexts (real BigMouth / KLE data) to pay off.

Table 2: Temperature scaling and ECE (synthetic-corpus baselines, 1,284 trials).

Model	T	ECE_{pre}	ECE_{post}
B0 frequency bigram	1.15	0.054	0.068
B1 TF-IDF + LR	1.15	0.054	0.031
B2 XGBoost	1.75	0.124	0.128

are already reasonably calibrated and only marginally improve.²

Applying the abstention rule “cover only when top-1 probability $\geq \tau=0.6$ ” yields the coverage-precision trade-off in Table 3:

The bigram baseline trades coverage for precision aggressively: it abstains on 79% of cases but answers correctly 77% of the time when it does. This is the pre-registered failure mode the LLM is expected to *mitigate* — matching that precision while abstaining on far fewer cases.

²B0’s ECE rises slightly after rescaling because the held-out half contains only 642 trials and 10-bin ECE has a noise floor of approximately 0.02 at this N ; the rescaling does not change ranking among models.




Figure 2: Top-1 accuracy by CDT category, apples-to-apples $n=30$ subsample. Classical baselines win the bundled / sequential categories (diagnostic, adjunctive, orthodontic) where the synthetic transition tables encode strong Markov structure. M2 and M3 *equal or beat* classical baselines on periodontic (0.667 vs ≤ 0.5) and preventive (1.0 for M2 vs ≤ 0.6). Restorative is hard for every model: chart-aware reasoning at $n=30$ is not enough signal to distinguish among single-surface, two-surface, and three-surface posterior composite codes that differ only in surface count.

Table 3: Coverage and precision under abstention at $\tau=0.6$.

Model	Coverage	Top-1 (overall)	Top-1 (covered)
B0 frequency bigram	0.21	0.556	0.772
B1 TF-IDF + LR	0.36	0.558	0.718
B2 XGBoost	0.51	0.507	0.654

4.4. Apples-to-apples LLM vs classical evaluation

We evaluate the LLM variants M1–M6 (defined in §3.3), all routed through the local Claude Code CLI per §3.6, on a 30-example random subsample of the test split (seed-deterministic, fixed across all three Sonnet+ models; M4 and M6 ran on a 15-example sub-subsample for compute reasons). Each LLM call is ~ 26 –40 seconds wall-clock through the local Claude Code CLI; total LLM runtime for this experiment was ~ 46 minutes for the 90-call M1+M2+M3 sweep, and a further ~ 25 minutes for M4 and M6 combined. For apples-to-apples comparison, we re-evaluate the four classical baselines (B0–B3) on the *same* 30 examples. Table 4 reports the




Figure 3: Reliability diagram for the three primary baseline models on the full $n=1,284$ baseline test split (LLMs omitted because $n=30$ is too small for stable per-bin estimates). The dashed diagonal is perfect calibration. B2 XGBoost is materially overconfident in the mid- range; B0 and B1 hew closer to the diagonal.

306 head-to-head.

Table 4: Apples-to-apples top-K accuracy on $n=30$ shared test examples. M4 evaluated on a $n=15$ subsample using the same seed; numbers are computed on the matching subset for direct comparison (where applicable).

Model	Top-1	Top-3	Top-5
B0 frequency bigram	0.567	0.800	0.933
B1 TF-IDF + LR	0.567	0.900	0.967
B2 XGBoost	0.433	0.933	0.967
B3 MultiTP-style CNN-RNN	0.567	0.867	0.900
M1 Claude Haiku	0.267	0.533	0.733
M2 Claude Sonnet + CoT	0.467	0.667	0.733
M3 Claude Sonnet + RAG	0.333	0.533	0.633
M4 Claude Opus + CoT ($n=15$)	0.467	0.667	0.733
M5 Sonnet + CoT + classical prior	0.533	0.767	0.933




Figure 4: Coverage-risk curve under top-1 abstention, $n=1,284$ baseline test split. As we lower the abstention threshold (cover more cases), risk rises. B2 XGBoost is dominated at every coverage level; B0 and B1 are within a few points of each other. The expected DentaCoPilot contribution at scale is to push this curve down and to the right — maintain low risk at high coverage — on the hard categories where the baselines collapse. LLM calibration cannot be read off this small synthetic-data run; it is a target for the BigMouth and KLE evaluations.

307 The classical baselines decisively beat the pure LLM variants (M1–M4)
 308 on this synthetic corpus. Three plausible reasons: (i) the synthetic tran-
 309 sition tables have very strong local Markov structure, which a frequency
 310 bigram captures perfectly with no chart-grounding required; (ii) the LLM
 311 is not given few-shot training examples in the prompt, so it must reason
 312 about a vocabulary of ~ 70 codes from descriptions alone; (iii) the LLMs
 313 frequently pick a category-correct code that misses the exact-CDT-match
 314 metric (e.g. D2391 where the gold is D2392 — both posterior composites
 315 differing only in surface count). The clinical interpretation of the latter
 316 failure mode is far less severe than top-1 exact-match suggests.

317 Two of the LLM-side findings are themselves results worth surfacing.
 318 First, the Sonnet $\rightarrow$ Opus comparison (M2 vs M4) shows *no improvement*
 319 from increased model capability: M4 scores identically to M2 (top-1 0.467,
 320 top-3 0.667, top-5 0.733). The Haiku $\rightarrow$ Sonnet step (M1 vs M2) does help
 321 (+0.20 top-1). The bottleneck on this task therefore is not raw model capa-

bility but prompt design, evaluation framing, and the synthetic data’s structure itself. Second, the M3 retrieval variant *underperforms* M2 — adding evidence cards to the prompt actively hurts on this corpus, suggesting the cards interact with the chain-of-thought in a way that distracts the model from the chart. Both are negative results in the sense of “does not improve top-1”; both are honest data points for the design space.

4.5. Prompt-conditioning on a classical prior closes the gap (M5)

The most important LLM-side result in our evaluation is M5. Injecting the top-10 (code, probability) candidates from the TF-IDF + logistic regression baseline into the Sonnet + chain-of-thought prompt — with explicit instruction that the model may re-rank, redistribute, or override the prior — changes the picture substantially:

- M5 top-1 (0.533) exceeds the best pure LLM (M2/M4 at 0.467) by +0.067 absolute, and approaches the classical baselines (B0/B1/B3 at 0.567).
- M5 top-3 (0.767) is +0.10 over M2.
- M5 top-5 (0.933) is +0.20 over M2 and matches B0 / B1 / B3 at top-5.

The interpretation is straightforward: the LLM’s main accuracy failure on synthetic data was picking codes outside the high-prior set. Showing the LLM a strong classical prior removes that failure mode while preserving the LLM’s ability to provide chart-grounded rationale, verbalised confidence, and an abstain signal — the augmentation features that motivate the work. M5 is the design we recommend going forward: a classical model handles the prior; the LLM provides the clinician-facing decision support layer on top.

Capability does not compound with priming.. An additional ablation, M6 (Opus + classical prior), tests whether the priming gain compounds with larger model capability. Apples-to-apples on the same $n=15$ keys:

Model	Top-1	Top-3	Top-5
M4 Opus alone	0.467	0.667	0.733
M5 Sonnet + B1 prior	0.467	0.733	0.933
M6 Opus + B1 prior	0.467	0.667	0.733

352 The prompt-conditioning gain that helps Sonnet (M2 $\rightarrow$ M5) does not
 353 translate to Opus (M4 $\rightarrow$ M6 is identity). Two plausible explanations:
 354 (i) Opus’s stronger built-in priors over-rule the injected classical prior more
 355 aggressively than Sonnet does; (ii) Opus’s chart reasoning already yields
 356 the relevant candidates without external help. The practical implication
 357 is favourable: **Sonnet + classical prior is both the best-performing**
 358 **and the cheapest LLM configuration we tested**. Sonnet inference is
 359 roughly 5 $\times$ cheaper per token than Opus, and the prior contributes the gain
 360 that capability alone does not.

361 4.5.1. Per-category breakdown

362 The headline numbers obscure a more interesting pattern visible in Fig-
 363 ure 2. On bundled / sequential procedures (diagnostic, adjunctive), classical
 364 baselines achieve ~ 0.89 top-1 while LLMs trail at ~ 0.33 – 0.47 . But on *pe-*
 365 *riodontic* and *preventive* procedures, M2 (Sonnet + CoT) achieves 0.667
 366 and 1.000 top-1 respectively, equalling or beating every classical baseline
 367 on those categories. M3 (Sonnet + RAG) shows the same pattern on pe-
 368 riodontic. The LLMs are not uniformly worse; they are differently worse,
 369 in a way that is consistent with chart-grounded reasoning helping on cases
 370 that require reasoning about the patient’s diabetic status, prior periodontal
 371 therapy, or recall cadence.

372 4.5.2. Verbalised confidence and abstention

373 Every LLM call emitted a confidence label in {low, medium, high} and
 374 an abstain flag. Across all three models, the abstain flag was never set
 375 (0/30 across M1/M2/M3); the prompt framing apparently does not push
 376 the models hard enough to use it, which we treat as a prompt-design issue
 377 rather than a model-capability claim. The empirical top-1 accuracies per
 378 confidence bucket:

Table 5: Empirical top-1 accuracy by self-reported confidence.

Model	low (n , acc)	medium (n , acc)	high (n , acc)
M1 Haiku	2, 0.000	12, 0.167	16, 0.375
M2 Sonnet + CoT	0, –	19, 0.526	11, 0.364
M3 Sonnet + RAG	0, –	22, 0.273	8, 0.500

379 Only M1 shows the expected monotonic relationship (low < medium <
 380 high). M2 is non-monotonic: “medium”-confidence predictions are more ac-

381 curate than “high”-confidence ones, which is the failure mode of verbalised
382 confidence as an LLM trustworthiness signal. M3 recovers monotonicity but
383 at lower overall accuracy. We interpret this as evidence that verbalised con-
384 fidence on its own is insufficient for abstention; calibrated top-1 probability
385 (which we do compute) is the better signal.

386 4.5.3. Sample rationale

387 An illustrative response from M2 on a periodontal-recall case (run_idm2_sonnet_cot-P000...-
388 verbatim from llm_calls.jsonl):

```
389     {"predictions": [{"code": "D0120", "probability": 0.38},  
390     {"code": "D1110", "probability": 0.28}, {"code": "D0274",  
391     "probability": 0.17}, {"code": "D2391", "probability": 0.10},  
392     {"code": "D2392", "probability": 0.07}], "confidence": "medium",  
393     "abstain": false, "rationale": "Last periodic evaluation  
394     was December 2022 and last prophylaxis visible in records  
395     was January 2021, making a routine recall visit (D0120 + D1110)  
396     the dominant next-step signal; annual bitewings (D0274) will  
397     also be due June 2024. The two recent posterior composites  
398     (teeth 17, 20) suggest ongoing caries activity, keeping one-  
399     and two-surface restorative codes (D2391, D2392) in the differential."}
```

400 The rationale identifies the two-and-a-half-year visit gap, recognises the
401 prior ~6–9-month recall cadence, prioritises the recall + prophylaxis + imaging
402 bundle, and keeps the surface-count restorative codes in the differential.
403 This is chart-grounded reasoning of a kind classical baselines categorically
404 cannot produce; whether this kind of reasoning *helps a dentist make a*
405 *better decision* is precisely what the forthcoming clinician-in-the-loop study
406 at KLE V K Institute is designed to measure, regardless of where exact-
407 CDT-match top-1 numbers land.

408 Critically, every LLM call in the system is routed through the local
409 Anthropic Claude Code CLI in preference to the metered Anthropic API,
410 and is logged in full to `data/results_<run>/llm_calls.jsonl` (prompt,
411 response, model identifier, latency, token counts). A reviewer may audit or
412 replay any decision the system makes without re-spending compute.

413 4.6. Confidence-routed gating: a negative result

414 Methods §3.5 described two ways to combine classical and LLM pre-
415 dictions; M5 above is the prompt-conditioning variant. We additionally
416 evaluated the confidence-routed-gating variant post-hoc: pick the baseline’s

top-1 when its calibrated probability $\geq \tau$, else the LLM’s top-1, sweeping $\tau \in \{0, 0.2, \dots, 1.0\}$ on the matched $n=30$ subsample for all four (B0–B3, M2) baseline-LLM pairings. At every τ , the hybrid’s top-1 accuracy was at most equal to the pure baseline’s; the gating did not produce a single threshold at which deferring to the LLM helped. We report this as a negative result: *when the LLM is uniformly weaker than the baseline on the target distribution, gating does not help*. The M5 prompt-conditioning result above is precisely the mechanism that does help in this regime — it lets the LLM use the baseline’s prior rather than *replace* it.

4.7. Real-data evaluation (planned)

Forthcoming work extends the M1–M6 evaluation to (a) the BigMouth multi-institutional dental data repository (Walji et al., 2014, 2022) pending controlled-access approval (over 4.5M patients across 11 academic contributing institutions in the 2022 update) and (b) a retrospective chart pull at the KLE V K Institute under an institutional ethics protocol. The clinician-in-the-loop study at KLE then evaluates dentist agreement, time-to-decision, override rate, and trust across three arms (dentist alone / dentist + classical recommender / dentist + DentaCoPilot), with all design parameters — sample size, primary and secondary outcomes, abstention threshold rule, stopping rules — locked in a pre-registration deposited on the Open Science Framework (Nosek et al., 2018) before any test-split unblinding.

5. Discussion

5.1. Hypotheses revisited

We posed two hypotheses in §1. The synthetic-corpus evidence is consistent with both: **(H1)** the M5 design (Sonnet + classical prior) ties classical baselines on top-5 accuracy (0.933) and narrows the top-1 gap to within four points (0.533 vs 0.567), while preserving chart-grounded rationale, verbalised confidence, and an abstain signal; **(H2)** prompt-conditioning is more cost-effective than upgrading the LLM backbone — M4 (Opus alone) and M6 (Opus + prior) match M2 (Sonnet alone) on every metric, so capability is not the lever once the prior is present. The findings stand on synthetic data only; whether they generalise to the BigMouth and KLE corpora is the empirical claim of the planned next stage.

450 5.2. *Augmentation, not autonomy*

451 The dominant framing for dental AI in 2026 is *augmented intelligence*
452 — AI as decision support, with the licensed clinician retaining diagnostic
453 and treatment authority. Our system design reflects that framing in three
454 concrete ways: a constrained vocabulary so the system cannot confidently
455 recommend nonsense codes; calibrated abstention so the system declines
456 instead of guesses; and a chart-grounded rationale so a dentist can audit
457 *why* a recommendation was made. The pre-registered clinician-in-the-loop
458 study at KLE specifically measures override rate alongside agreement; we
459 treat the dentist’s freedom to disagree with the model as a feature, not a
460 calibration problem.

461 5.3. *Reproducibility as a first-class outcome*

462 Every experimental decision in this work is recorded such that a reviewer
463 with access to a Claude Code subscription can replay the entire pipeline at
464 zero marginal cost. Each LLM call is logged with prompt, response, model
465 identifier, latency, and token counts; per-run `config.json` files record the
466 git SHA, seed, and CLI arguments. The synthetic chart generator and
467 the BM25 retriever over the evidence corpus are open-sourced under MIT
468 licence. The pre-registered analysis plan is deposited on OSF before test-
469 split unblinding. No part of the synthetic-data results or the calibration /
470 abstention analyses requires the BigMouth or KLE clinical data: these are,
471 deliberately, the parts a reviewer can verify today.

472 5.4. *Limitations*

473 *Synthetic-data ceiling.* The current evaluation runs on plausibility-based
474 synthetic charts. Real dental practice contains patterns we did not encode
475 (insurance-driven sequencing, specialist referral cycles, multi-visit treatment
476 phases, idiosyncratic per-clinic preferences) and codes we omit from the rep-
477 resentative vocabulary. The synthetic results validate the pipeline; they do
478 not validate the system. The BigMouth and KLE experiments are required
479 for any substantive accuracy claim.

480 *Sample sizes.* The reported LLM evaluation is on 30 test examples per
481 model. Statistical comparisons against the baselines are under-powered; we
482 report point estimates and avoid significance claims. A 200-example follow-
483 up evaluation is straightforward and is the next planned run.

484 *Vocabulary coverage.* We constrain LLM outputs to a representative subset
485 of ~ 70 frequent CDT 2024 codes plus a training-set-derived extension. Rare
486 codes are absent. The full CDT codeset (>700 codes) is ADA copyright; a
487 downstream user must merge the official codeset before clinical deployment.

488 *Calibration is local.* Temperature scaling is fit on a random half of the test
489 split, which is a within-corpus calibration. It does not address distribution
490 shift between dental schools or between the synthetic and real settings. We
491 expect to refit calibration on each target corpus.

492 *Abstention was not used.* Across all 90 LLM trials in the $n=30$ evaluation,
493 none of M1, M2, or M3 emitted `abstain: true`. The prompt instructed
494 the model to abstain when “chart context is genuinely insufficient,” but
495 the synthetic corpus rarely produces such inputs, and the prompt does not
496 provide concrete examples of abstainable cases. This is a genuine failure of
497 our prompt design — not a capability claim about the models. Future work
498 will explicitly construct adversarial / out-of-distribution chart inputs and
499 either include them in few-shot exemplars or replace the verbalised abstain
500 flag with a learned threshold over calibrated top-1 probability.

501 5.5. Ethics and clinical safety

502 We deliberately do not seek FDA or other regulatory clearance for this
503 system. We do not deploy it. The pre-registered clinician-in-the-loop study
504 at the KLE V K Institute will be conducted under the institution’s stan-
505 dard ethics-review process; no clinician evaluates a single case before that
506 approval is in place. The retrospective chart pull will use only fully de-
507 identified records. Patient-identifiable data never leaves the institutional
508 firewall; all large-language-model inference runs locally through Claude
509 Code, with no patient identifiers transmitted to any third-party LLM ser-
510 vice. Author contributions follow the ICMJE recommendations (Interna-
511 tional Committee of Medical Journal Editors, 2024).

512 *Conflicts of interest.* The authors declare no financial relationship with
513 Anthropic, Overjet, Pearl, VideaHealth, or any other commercial dental-AI
514 vendor. CELABE purchases a Claude Code subscription at standard retail
515 terms; Anthropic provided no funding, no preferred access, and no editorial
516 input on this manuscript. C.C.R. is the founder of CELABE; CELABE
517 has not filed a patent and has no commercial product relating to this work.
518 S.D.R. is a final-year BDS candidate at the KLE V K Institute and declares

no competing interests. We invite scrutiny of both the design and the implementation.

5.6. Future work

Three directions: (i) the BigMouth and KLE experiments described in the pre-registration; (ii) retrieval over a properly licensed clinical- guideline corpus when one becomes available, and ablation of M2 vs M3; and (iii) paired imaging context for the most common dental decisions — incorporating panoramic / periapical findings (DENTEX, DenPAR) as additional structured chart context for the LLM, rather than as a standalone classification target.

6. Conclusion

We presented the first stage of *DentaCoPilot*, an LLM-augmented next-procedure recommender for general dentistry that is designed, prompted, and evaluated as a decision-support consult instead of an autonomous classifier. On a synthetic chart corpus, we found that **(a)** pure LLM variants (Haiku, Sonnet + chain-of-thought, Sonnet + retrieval, Opus + chain-of-thought) trail classical baselines on top-1 exact-CDT-match accuracy by 10–30 percentage points; **(b)** *prompt-conditioning a Sonnet LLM on a classical baseline’s top-10 candidates (M5) closes the gap*, raising top-5 from 0.733 to 0.933, matching classical baselines, while preserving chart-grounded rationale, verbalised confidence, and an abstain signal; **(c)** upgrading the LLM backbone from Sonnet to Opus does not improve performance with or without priming, indicating that capability is not the lever once the prior is present. The M5 design is, on this evidence, both the best-performing and the cheapest LLM configuration we tested. Every LLM call in the system is logged for full audit and replayable at zero marginal cost on a Claude Code subscription. The pre-registered next stage of work — the multi-institutional BigMouth evaluation and the dentist-in-the-loop study at the KLE V K Institute of Dental Sciences — will be the empirical test of whether the M5 design generalises beyond synthetic data and whether it produces a tool dentists actually use.

Acknowledgements

The authors acknowledge the use of Anthropic Claude (Claude Code) for code authoring and manuscript drafting assistance during this work. All

553 clinical and methodological decisions, as well as final interpretation, were
554 made by the human authors.

555 **AI Declaration**

556 This manuscript was prepared with assistance from Anthropic Claude
557 (Claude Code). Specifically, Claude was used for code scaffolding, draft
558 prose authoring, and reference management. The authors verified every
559 claim, edited all generated text, and take full responsibility for the content.
560 The model is also a direct subject of this work: all experimental large
561 language model calls in the system under evaluation were routed through
562 Claude Code rather than the Anthropic API, an explicit choice documented
563 in the methods section.

564 **References**

- 565 American Dental Association, 2024. CDT 2024: Current dental terminology Procedural
566 code set used throughout this work.
- 567 Anthropic, 2026. Claude code: an agentic coding environment <https://docs.claude.com/claude-code> – compute substrate for this paper’s experiments.
- 568 Chen, H., Liu, P., Chen, Z., Chen, Q., Wen, Z., Xie, Z., 2024. Predicting sequenced
569 dental treatment plans from electronic dental records using deep learning. *Artificial*
570 *Intelligence in Medicine* 148, 102734. doi:10.1016/j.artmed.2023.102734.
- 571 Chen, T., Guestrin, C., 2016. XGBoost: a scalable tree boosting system, in: *KDD*.
- 572 El-Yaniv, R., Wiener, Y., 2010. Foundations of selective prediction. *Journal of Machine*
573 *Learning Research* 11, 1605–1641.
- 574 Farhadi Nia, M., Ahmadi, M., Irankhah, E., 2024. Transforming dental diagnostics with
575 artificial intelligence: advanced integration of ChatGPT and large language models
576 for patient care. *Frontiers in Dental Medicine* 5, 1456208. doi:10.3389/fdmed.2024.
577 1456208.
- 578 Geifman, Y., El-Yaniv, R., 2017. Selective classification for deep neural networks, in:
579 *NeurIPS*.
- 580 Guo, C., Pleiss, G., Sun, Y., Weinberger, K.Q., 2017. On calibration of modern neural
581 networks .
- 582 Hamamci, I.E., Er, S., Simsar, E., Yuksel, A.E., Gultekin, S., Ozdemir, S.D., Yang,
583 K., Li, H.B., Pati, S., Stadlinger, B., et al., 2023. DENTEX: a multi-institutional
584 panoramic dental X-ray challenge. *MICCAI 2023 Grand Challenge*; <https://dentex.grand-challenge.org/>.
- 585 International Committee of Medical Journal Editors, 2024. Recommendations for the
586 conduct, reporting, editing, and publication of scholarly work in medical journals
587 <https://www.icmje.org/recommendations/>.
- 588 Kadavath, S., Conerly, T., Askell, A., et al., 2022. Language models (mostly) know what
589 they know. *arXiv:2207.05221* .

592 Kalenderian, E., Walji, M.F., et al., 2018. Dental procedure-code use in the united states:
 593 an analysis of variability. *Journal of the American Dental Association* .
 594 Li, Z., Wang, L., 2025. Multi-task reinforcement learning and explainable
 595 AI-driven platform for personalized planning and clinical decision support in
 596 orthodontic-orthognathic treatment. *Scientific Reports* 15, 24502. doi:10.1038/
 597 s41598-025-09236-z.
 598 Lin, S., Hilton, J., Evans, O., 2022. Teaching models to express their uncertainty in
 599 words. *Transactions on Machine Learning Research* .
 600 Mertens, S., Krois, J., Cantu, A.G., Arsiwala, L.T., Schwendicke, F., 2021. Artificial
 601 intelligence for caries detection: randomized trial. *Journal of Dentistry* 115, 103849.
 602 doi:10.1016/j.jdent.2021.103849.
 603 Munafò, M.R., Nosek, B.A., Bishop, D.V.M., et al., 2017. A manifesto for reproducible
 604 science. *Nature Human Behaviour* 1, 0021.
 605 Naeini, M.P., Cooper, G.F., Hauskrecht, M., 2015. Obtaining well calibrated probabilities
 606 using Bayesian binning, in: *AAAI*.
 607 Nixon, J., Dusenberry, M.W., Zhang, L., Jerfel, G., Tran, D., 2019. Measuring calibration
 608 in deep learning, in: *CVPR Workshops*.
 609 Nori, H., King, N., McKinney, S.M., et al., 2023. Capabilities of GPT-4 on medical
 610 challenge problems. *arXiv:2303.13375* .
 611 Nosek, B.A., Ebersole, C.R., DeHaven, A.C., Mellor, D.T., 2018. The preregistration
 612 revolution. *Proceedings of the National Academy of Sciences* 115, 2600–2606.
 613 Omiye, J.A., Lester, J.C., Spichak, S., Rotemberg, V., Daneshjou, R., 2024. Large
 614 language models propagate race-based medicine. *npj Digital Medicine* .
 615 Park, Y., et al., 2024. Methodological standards for clinical machine learning research.
 616 *npj Digital Medicine* .
 617 Patel, J., Badi, M., Katiyar, R., Ogwo, C., Wiener, R.C., Tiwari, T., Sambamoorthi,
 618 U., Folks, T., 2025. SDoH impact on periodontal disease using machine learning and
 619 dental records. *Journal of Dental Research* doi:10.1177/00220345251328968.
 620 Patel, J.S., Tellez, M., Katiyar, R., Al-Hebshi, N.N., Santana, R., Yucel, R.M., Ismail, A.,
 621 2026. Periodontitis prediction model using linked electronic health and dental records.
 622 *JDR Clinical & Translational Research* doi:10.1177/23800844251408849. online first;
 623 print 2026.
 624 Pedregosa, F., Varoquaux, G., Gramfort, A., et al., 2011. Scikit-learn: machine learning
 625 in Python. *Journal of Machine Learning Research* 12, 2825–2830.
 626 Platt, J.C., 2000. Probabilistic outputs for support vector machines and comparisons to
 627 regularized likelihood methods. *Advances in Large Margin Classifiers* .
 628 Rajkomar, A., Oren, E., Chen, K., Dai, A.M., et al., 2018. Scalable and accurate deep
 629 learning with electronic health records. *npj Digital Medicine* .
 630 Rasnayaka, S., Bandara, D.L., Jayasundara, A., Jayasinghe, R., Wimalasiri, C., Rath-
 631 nayake, P., Wijerathne, S., Ragel, R., Thambawita, V., Nawinne, I., 2025. DenPAR:
 632 annotated intra-oral periapical radiographs dataset for machine learning. *Scientific*
 633 *Data* 12, 1615. doi:10.1038/s41597-025-05906-9.
 634 Robertson, S.E., Walker, S., Jones, S., Hancock-Beaulieu, M.M., Gatford, M., 1995.
 635 Okapi at TREC-3, in: *Proceedings of the Third Text REtrieval Conference*.
 636 Salton, G., Buckley, C., 1988. Term-weighting approaches in automatic text retrieval.
 637 *Information Processing & Management* 24, 513–523.

638 Schwendicke, F., Mertens, S., Cantu, A.G., Chaurasia, A., Meyer-Lueckel, H., Krois,
 639 J., 2022. Cost-effectiveness of AI for caries detection: randomized trial. *Journal of*
 640 *Dentistry* 119, 104080. doi:10.1016/j.jdent.2022.104080.
 641 Schwendicke, F., Samek, W., Krois, J., 2020. Artificial intelligence in dentistry:
 642 chances and challenges. *Journal of Dental Research* 99, 769–774. doi:10.1177/
 643 0022034520915714.
 644 Seifo, N., Cassie, H., Radford, J.R., Innes, N.P.T., 2020. Silver diamine fluoride for caries
 645 arrest in primary teeth: a systematic review. *British Dental Journal* .
 646 Shujaat, S., 2026. FDA-approved AI solutions in dental imaging: a narrative review
 647 of applications, evidence, and outlook. *International Dental Journal* doi:10.1016/j.
 648 identj.2025.109315. online December 2025.
 649 Singhal, K., Azizi, S., Tu, T., et al., 2023. Large language models encode clinical knowl-
 650 edge. *Nature* 620, 172–180.
 651 Slayton, R.L., Urquhart, O., Araujo, M.W.B., et al., 2018. Evidence-based clinical
 652 practice guideline on nonrestorative treatments for carious lesions. *Journal of the*
 653 *American Dental Association* .
 654 Varghese, J., Sankeshwari, R., Nagamoti, J.M., Jigan, J., et al., 2025. Cariogram-based
 655 caries risk profile in preschoolers with parent-oriented educational mobile messages.
 656 *International Journal of Clinical Pediatric Dentistry* .
 657 Wadia, R., 2017. Tooth-numbering systems. *British Dental Journal* 223, 84.
 658 Walji, M.F., Kalenderian, E., Stark, P.C., White, J.M., Kookal, K.K., Phan, D., Tran,
 659 D., Bernstam, E.V., Ramoni, R., 2014. BigMouth: a multi-institutional dental data
 660 repository. *Journal of the American Medical Informatics Association* 21, 1136–1140.
 661 doi:10.1136/amiajnl-2013-002230.
 662 Walji, M.F., Spallek, H., Kookal, K.K., Barrow, J., Magnuson, B., Tiwari, T., Oy-
 663 oyo, U., Brandt, M., Howe, B.J., Anderson, G.C., White, J.M., Kalenderian,
 664 E., 2022. BigMouth: development and maintenance of a successful dental data
 665 repository. *Journal of the American Medical Informatics Association* 29, 701–706.
 666 doi:10.1093/jamia/ocab281.
 667 Wang, C., Zhang, Y., Wu, C., Liu, J., Huang, X., Wu, L., Wang, Y., Feng, X., Lu, Y.,
 668 Wang, Y., 2025. MMDental: a multimodal dataset of tooth CBCT images with expert
 669 medical records. *Scientific Data* 12, 1172. doi:10.1038/s41597-025-05398-7.
 670 Wei, J., Wang, X., Schuurmans, D., Bosma, M., Chi, E., Le, Q., Zhou, D., 2022. Chain-
 671 of-thought prompting elicits reasoning in large language models. *NeurIPS* .
 672 Xiong, M., Hu, Z., Lu, X., et al., 2024. Can LLMs express their uncertainty?, in: *ICLR*.